

COMP 333 I/933 I: Computer Networks and Applications

Week 6

Congestion Control (Transport Layer)

Reading Guide: Chapter 3, Sections: 3.6-3.7

Transport Layer: Outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control


3.7 TCP congestion control

Principles of congestion control

congestion:


- ❖ informally: “too many sources sending too much data too fast for *network* to handle”
- ❖ different from flow control!
- ❖ manifestations:
 - lost packets (buffer overflow at routers)
 - long delays (queueing in router buffers)
- ❖ a top-10 problem!

Congestion



Incoming rate is faster than
outgoing link can support.

Congestion



Incoming rate is faster than
outgoing link can support.

Quiz: What's the worst that can happen?




A: This is no problem. Senders just keep transmitting, and it'll all work out.

B: There will be retransmissions, but the network will still perform without much trouble.

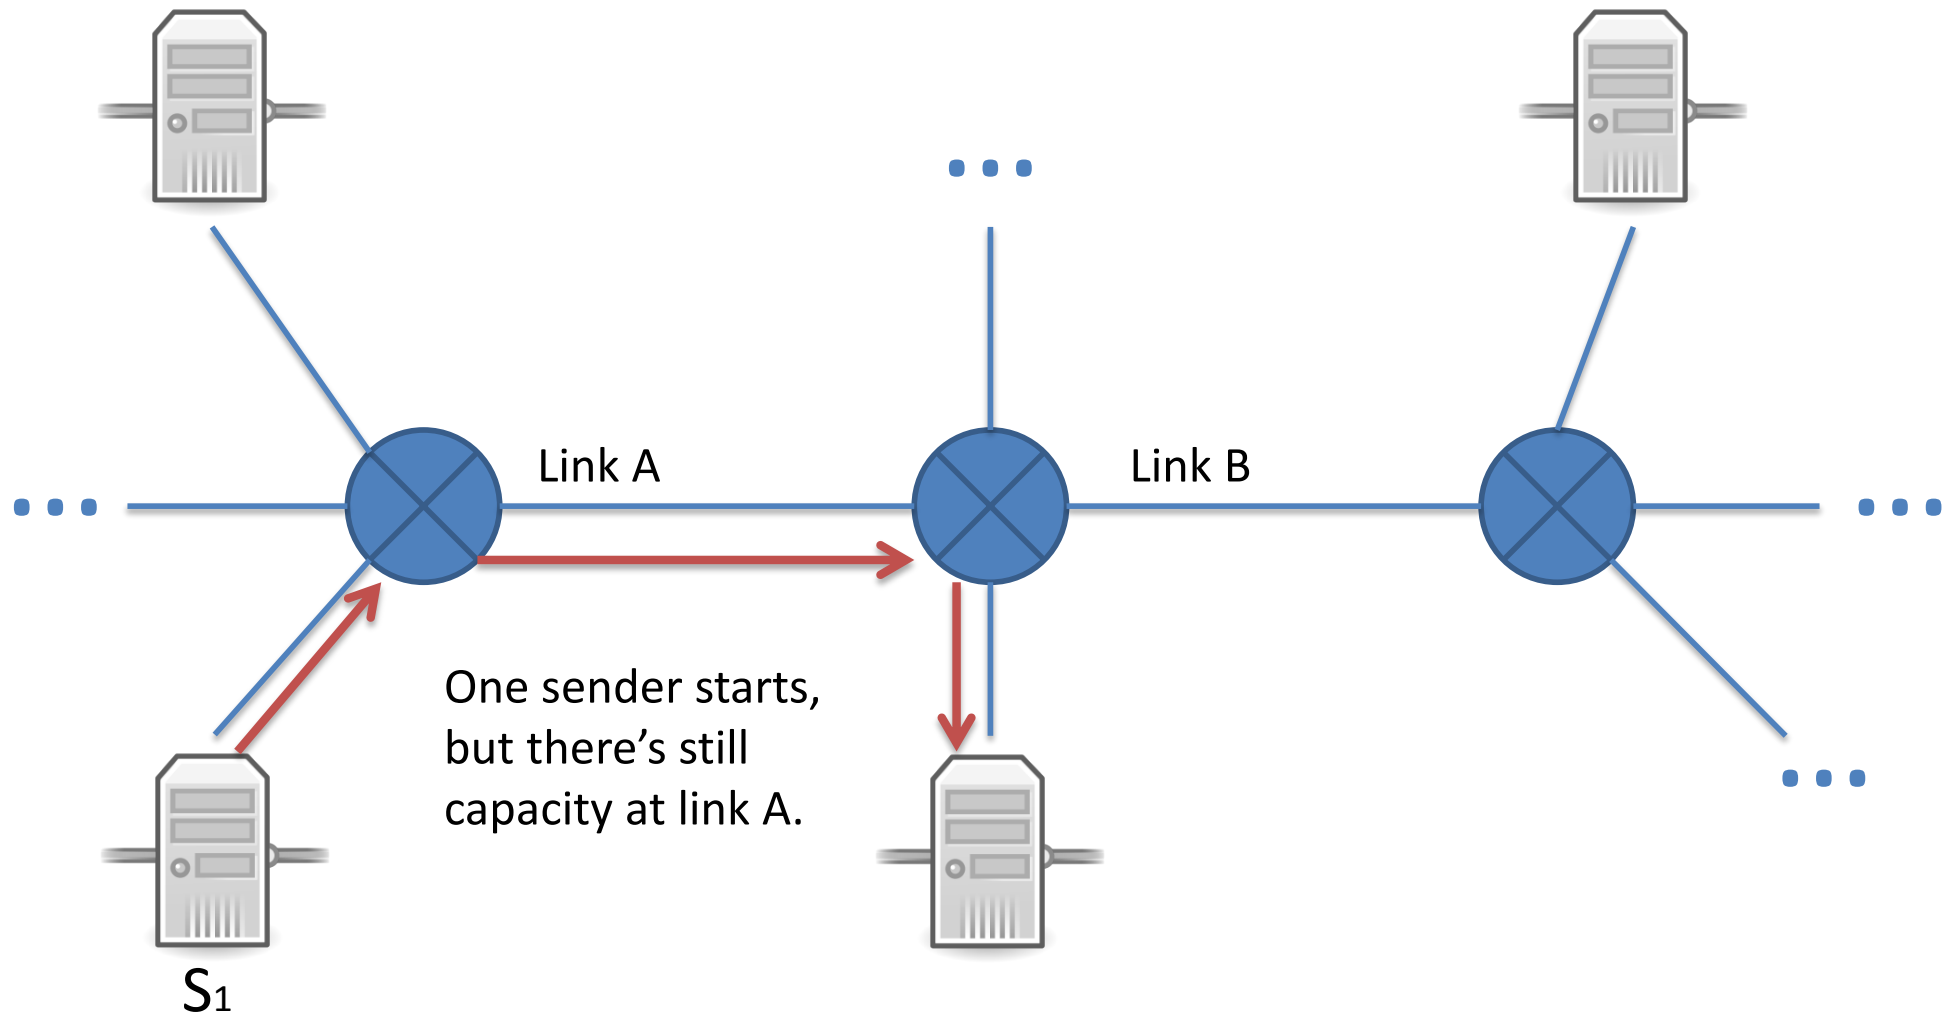
C: Retransmissions will become very frequent, causing a serious loss of efficiency

D: The network will become completely unusable

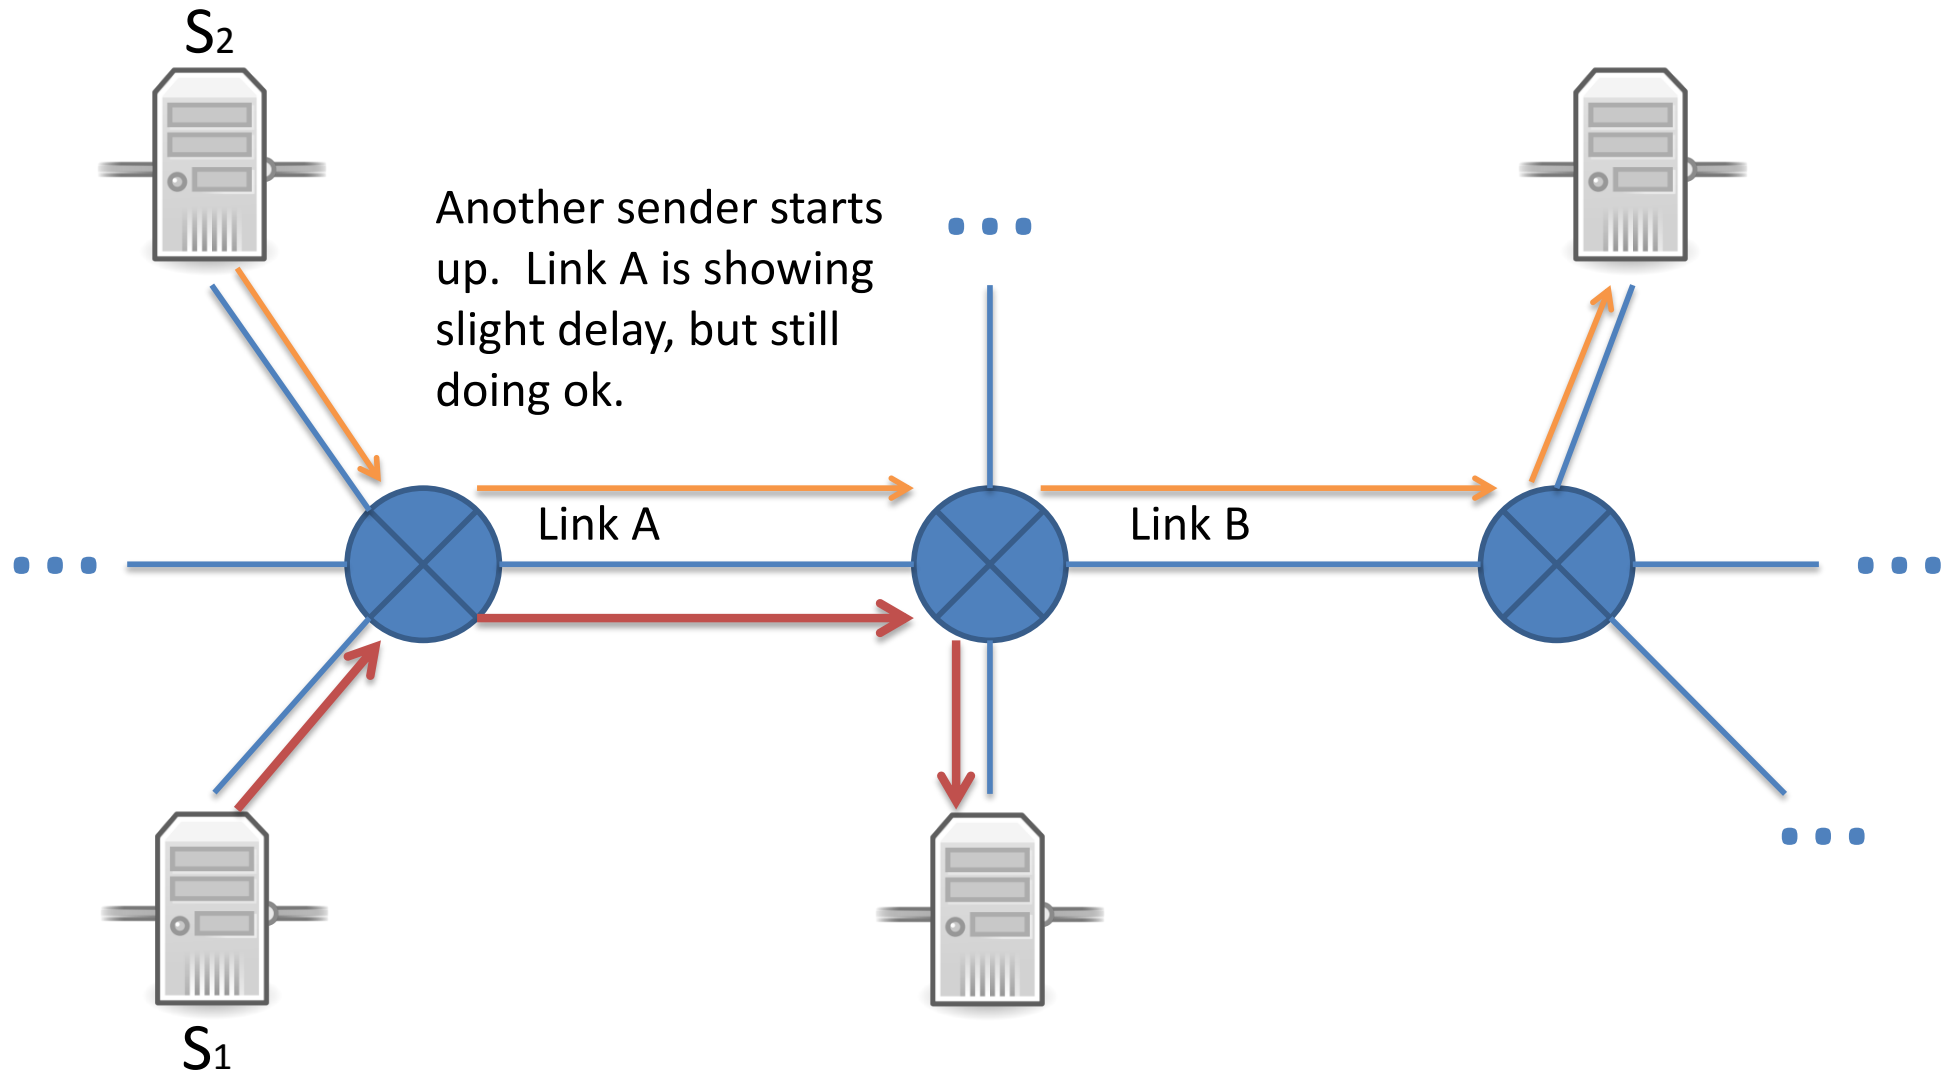
Congestion Collapse



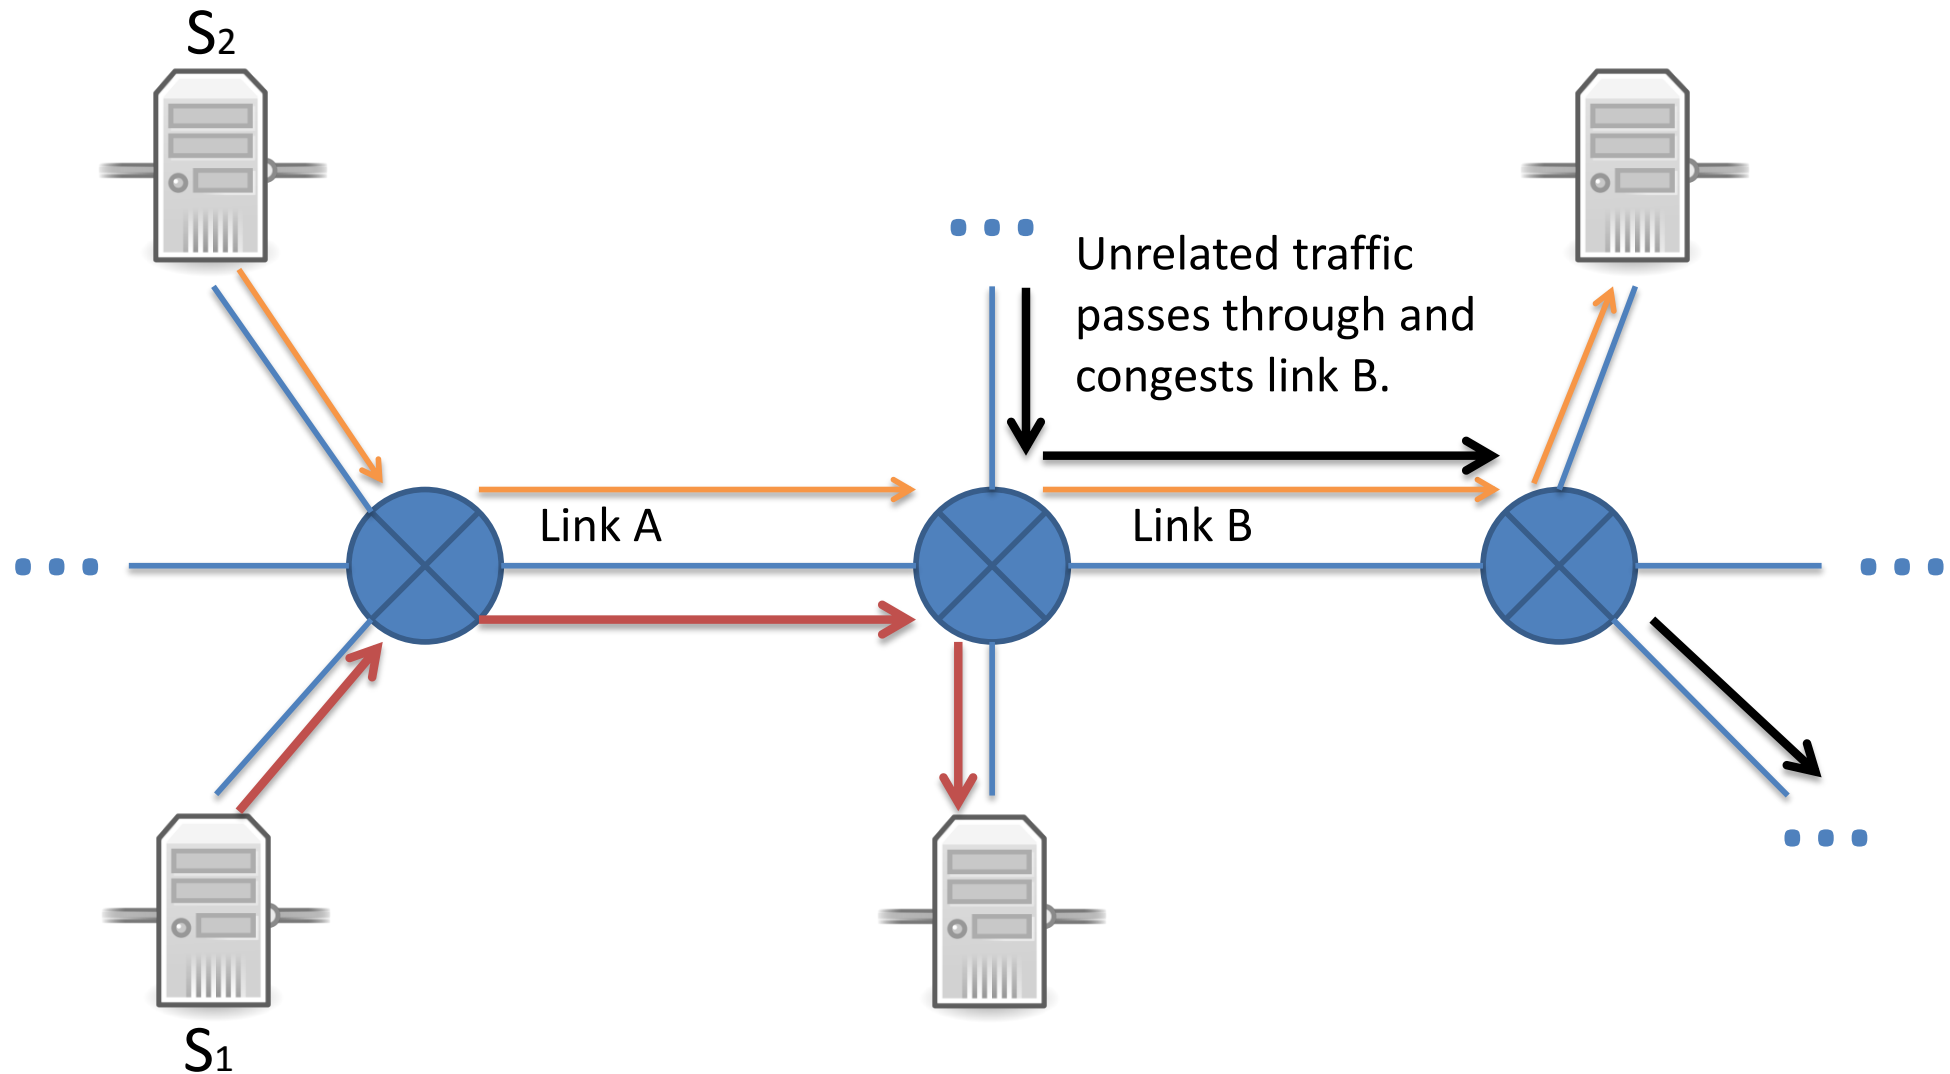
Congestion Collapse



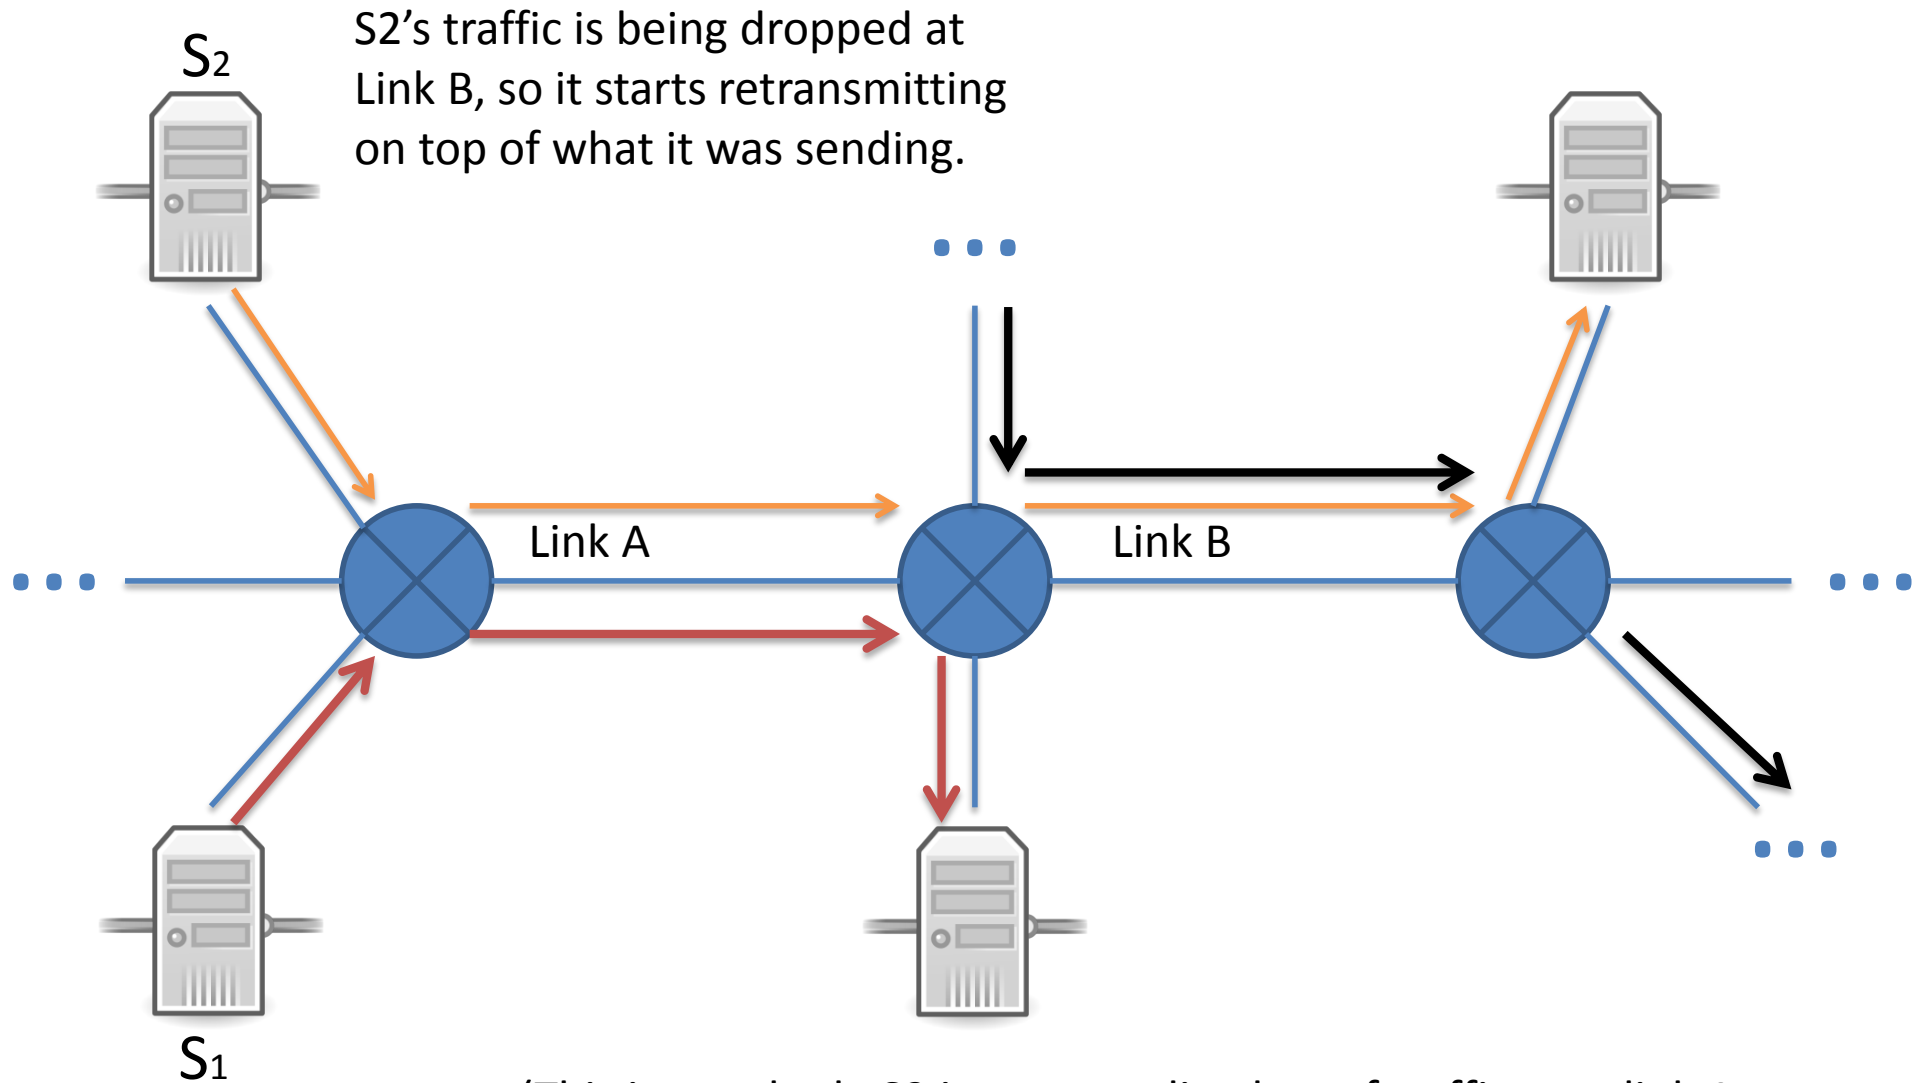
Congestion Collapse



Congestion Collapse

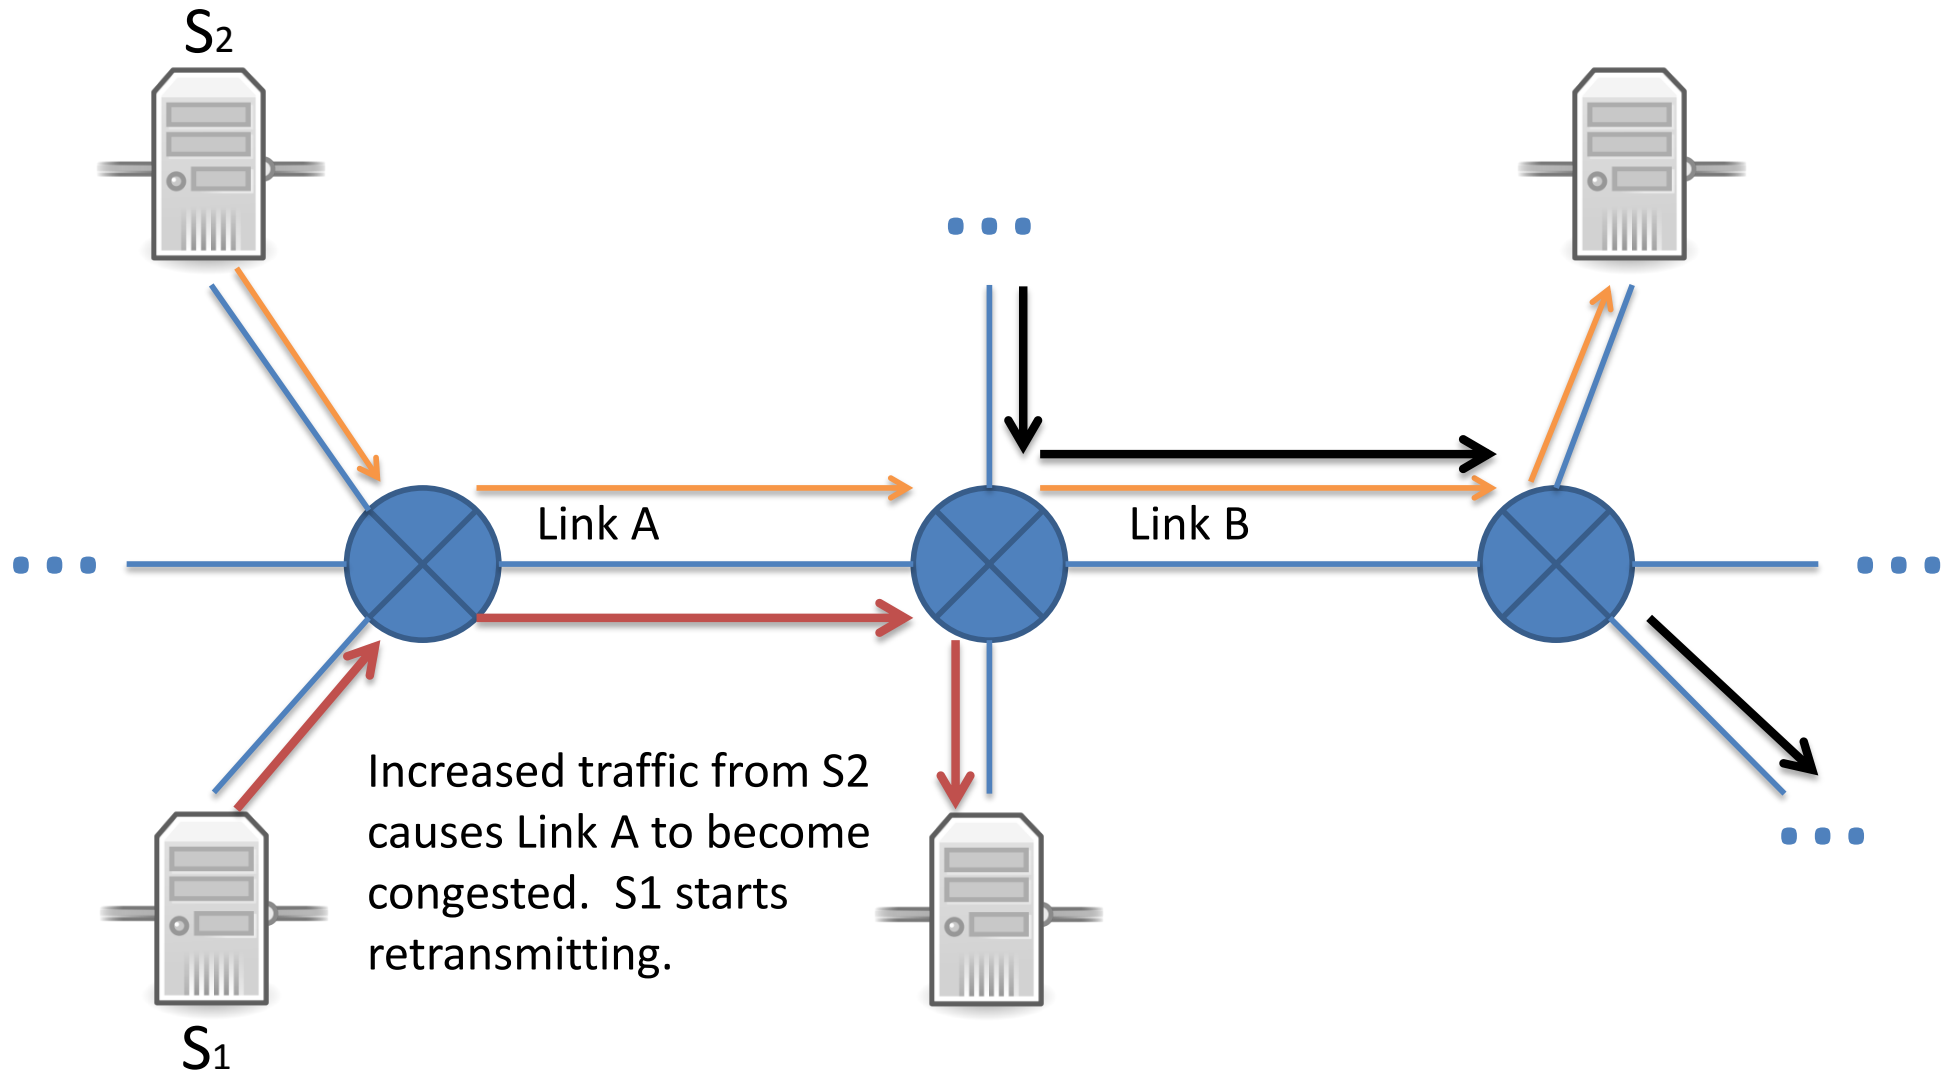


Congestion Collapse




(This is very bad. S₂ is now sending lots of traffic over link A that has no hope of crossing link B.)

Congestion Collapse



Congestion Collapse




Without congestion control

congestion:

- ❖ Increases delivery latency
- ❖ Increases loss rate
- ❖ Increases retransmissions, many unnecessary
- ❖ Wastes capacity of traffic that is never delivered
- ❖ Increases congestion, cycle continues ...

Cost of Congestion

- ❖ Knee – point after which
 - Throughput increases slowly
 - Delay increases fast
- ❖ Cliff – point after which
 - Throughput starts to drop to zero (congestion collapse)
 - Delay approaches infinity



Congestion Collapse

This happened to the Internet (then NSFnet) in 1986

- ❖ Rate dropped from a blazing 32 Kbps to 40bps
- ❖ This happened on and off for *two years*
- ❖ In 1988, Van Jacobson published “Congestion Avoidance and Control”
- ❖ The fix: senders voluntarily limit sending rate

Approaches towards congestion control

two broad approaches towards congestion control:

end-end congestion control:

- ❖ no explicit feedback from network
- ❖ congestion inferred from end-system observed loss, delay
- ❖ approach taken by TCP

network-assisted congestion control:

- ❖ routers provide feedback to end systems
 - single bit indicating congestion (SNA, DECbit, TCP/IP ECN, ATM)
 - explicit rate for sender to send at

Transport Layer: Outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management


3.6 principles of congestion control

3.7 TCP congestion control

TCP's Approach in a Nutshell

- ❖ TCP connection has window
 - Controls number of packets in flight
- ❖ *TCP sending rate:*
 - *roughly:* send cwnd bytes, wait RTT for ACKS, then send more bytes

$$\text{rate} \approx \frac{\text{cwnd}}{\text{RTT}} \text{ bytes/sec}$$



- ❖ **Vary window size to control sending rate**

All These Windows...

- ❖ Congestion Window: **CWND**
 - How many bytes can be sent without overflowing routers
 - Computed by the sender using congestion control algorithm
- ❖ Flow control window: **AdvertisedWindow (RWND)**
 - How many bytes can be sent without overflowing receiver's buffers
 - Determined by the receiver and reported to the sender
- ❖ Sender-side window = **minimum**{**CWND**, **RWND**}
 - Assume for this lecture that $RWND \gg CWND$

CWND

- ❖ This lecture will talk about CWND in units of MSS
 - (Recall MSS: Maximum Segment Size, the amount of payload data in a TCP packet)
 - This is only for pedagogical purposes
- ❖ Keep in mind that real implementations maintain CWND in bytes

Two Basic Questions

- ❖ How does the sender detect congestion?
- ❖ How does the sender adjust its sending rate?

Quiz: What is a “congestion event”



A: A segment loss (but how can the sender be sure of this?)

B: Increased delays

C: Receiving duplicate acknowledgement (s)

D: A retransmission timeout firing

D: Some subset of A, B, C & D (what is the subset?)

Quiz: How should we set CWND?



A: We should keep raising it until a “congestion event” then back off slightly until we notice no more events.

B: We should raise it until a “congestion event”, then go back to 1 and start raising it again

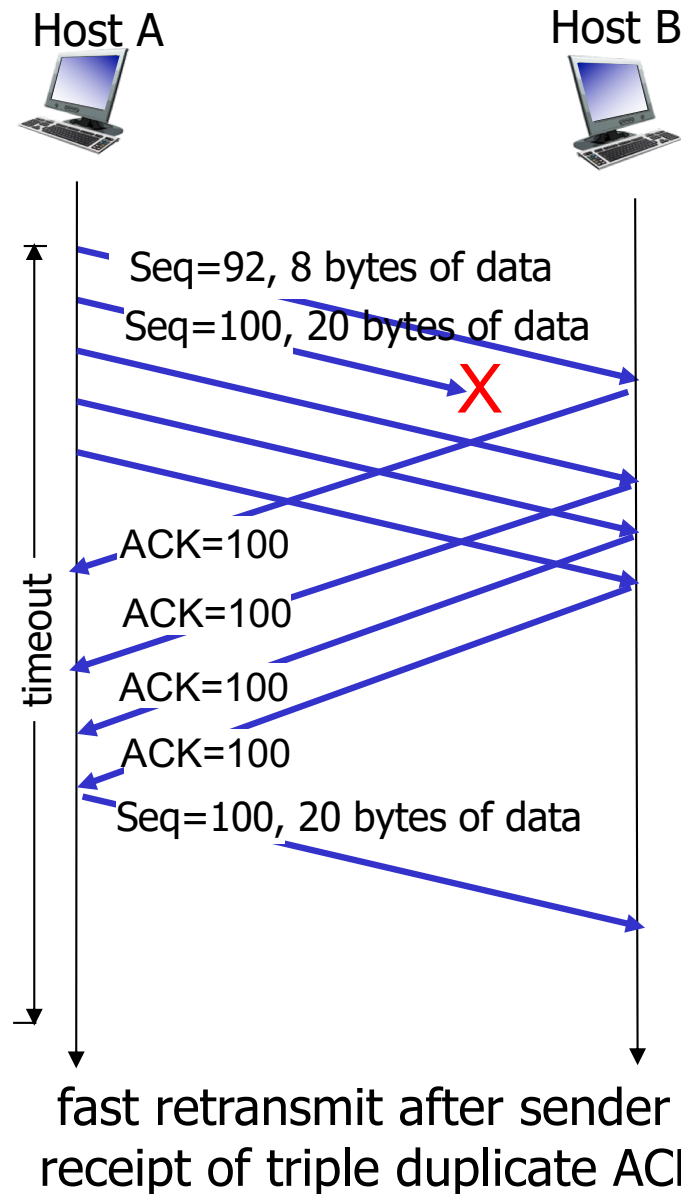
C: We should raise it until a “congestion event” then go back to median value and start raising it again.

D: We should sent as fast as possible at all times.

Not All Losses the Same

- ❖ Duplicate ACKs: isolated loss
 - dup ACKs indicate network capable of delivering some segments
- ❖ Timeout: much more serious
 - Not enough dup ACKs
 - Must have suffered several losses
- ❖ Will adjust rate differently for each case

RECAP: TCP fast retransmit



Rate Adjustment


- ❖ Basic structure:
 - Upon receipt of ACK (of new data): increase rate
 - Upon detection of loss: decrease rate
- ❖ How we increase/decrease the rate depends on the phase of congestion control we're in:
 - Discovering available bottleneck bandwidth vs.
 - Adjusting to bandwidth variations

Bandwidth Discovery with Slow Start (SS)

- ❖ Goal: estimate available bandwidth
 - start slow (for safety)
 - but ramp up quickly (for efficiency)
- ❖ Consider
 - $RTT = 100\text{ms}$, $MSS = 1000\text{bytes}$
 - Window size to fill 1Mbps of BW = 12.5 packets
 - Window size to fill 1Gbps = 12,500 packets
 - Either is possible!

TCP Slow Start

- ❖ when connection begins, increase rate exponentially until first loss event:
 - initially `cwnd` = 1 MSS
 - double `cwnd` every RTT
 - Simpler implementation achieved by incrementing `cwnd` for every ACK received
- ❖ summary: initial rate is slow but ramps up exponentially fast




Adjusting to Varying Bandwidth

- ❖ Slow start gave an estimate of available bandwidth
- ❖ Now, want to track variations in this available bandwidth, oscillating around its current value
 - Repeated probing (rate increase) and backoff (rate decrease)
 - Known as Congestion Avoidance (CA)
- ❖ TCP uses: “Additive Increase Multiplicative Decrease” (AIMD)
 - We’ll see why shortly...


AIMD

- ❖ *approach*: sender increases transmission rate (window size), probing for usable bandwidth, until loss occurs
 - *additive increase*: increase **cwnd** by 1 MSS every RTT until loss detected
 - For each successful RTT, $\text{cwnd} = \text{cwnd} + 1$
 - Simple implementation: for each ACK, $\text{cwnd} = \text{cwnd} + 1/\text{cwnd}$
 - *multiplicative decrease*: cut **cwnd** in half after loss

AIMD saw tooth behavior: probing for bandwidth



Leads to the TCP “Sawtooth”



Slow-Start vs. AIMD

- ❖ When does a sender stop Slow-Start and start Additive Increase?
- ❖ Introduce a “slow start threshold” (**ssthresh**)
 - Initialized to a large value
 - On timeout/loss, $ssthresh = CWND/2$
- ❖ When $CWND = ssthresh$, sender switches from slow-start to AIMD-style increase

Implementation

❖ State at sender

- **CWND** (initialized to a small constant)
- **ssthresh** (initialized to a large constant, so initial slow start can learn network condition fast)
- [Also **dupACKcount** and **timer**, as before]

❖ Events

- ACK (new data)
- dupACK (duplicate ACK for old data)
- Timeout

Event: ACK (new data)

❖ If $CWND < ssthresh$

■ $CWND += 1$

- $CWND$ packets per RTT
- Hence after one RTT with no drops:
 $CWND = 2 \times CWND$

Event: ACK (new data)

- ❖ If $CWND < ssthresh$
 - $CWND += 1$

Slow start phase

- ❖ Else
 - $CWND = CWND + 1/CWND$

*“Congestion Avoidance” phase
(additive increase)*

- $CWND$ packets per RTT
- Hence after one RTT
with no drops:
 $CWND = CWND + 1$

Event: dupACK


- ❖ dupACKcount ++
- ❖ If dupACKcount = 3 /* fast retransmit */
 - ssthresh = CWND/2
 - **CWND = CWND/2**

Event: Timeout

❖ On Timeout


- $\text{ssthresh} \leftarrow \text{CWND}/2$
- $\text{CWND} \leftarrow 1$

Example




Slow-start restart: Go back to $CWND = 1$ MSS, but take advantage of knowing the previous value of $CWND$

TCP Flavours



- ❖ TCP-Reno (Assumed Default in this course)
 - $\text{cwnd} = 1$ on timeout
 - $\text{cwnd} = \text{cwnd}/2$ on triple dup ACK
- ❖ TCP-Tahoe (Old/original version)
 - $\text{cwnd} = 1$ on both triple dup ACK & timeout
- ❖ Figure 3.52, page 304 of 7th Ed. textbook assumes a special TCP Reno that implements Fast Recovery, which is out of scope in this course.

TCP/Reno (Default in this course): Big Picture



TD: Triple duplicate acknowledgements
TO: Timeout

Transport Layer: Summary

- ❖ principles behind transport layer services:
 - multiplexing, demultiplexing
 - reliable data transfer
 - flow control
 - congestion control
- ❖ instantiation, implementation in the Internet
 - UDP
 - TCP

next:

- ❖ leaving the network “edge” (application, transport layers)
- ❖ into the network “core”